

Forest Cover Classification

DT Nicolay 26296918
Department of Computer Science
Stellenbosch University
Stellenbosch, South Africa
26296918@sun.ac.za

Abstract—This document is a model and instructions for L^AT_EX. This and the IEEEtran.cls file define the components of your paper [title, text, heads, etc.]. *CRITICAL: Do Not Use Symbols, Special Characters, Footnotes, or Math in Paper Title or Abstract.

Index Terms—component, formatting, style, styling, insert.

I. INTRODUCTION

introooo TODO remove “Fig. ??”, even at the beginning of a sentence.

TABLE I
TABLE TYPE STYLES

| Table Head | Table Column Head | | |
|------------|------------------------------|---------|---------|
| | Table column subhead | Subhead | Subhead |
| copy | More table copy ^a | | |

^aSample of a Table footnote.

Figure Labels: Use 8 point Times New Roman for Figure labels. Use words rather than symbols or abbreviations when writing Figure axis labels to avoid confusing the reader. As an example, write the quantity “Magnetization”, or “Magnetization, M”, not just “M”. If including units in the label, present them within parentheses. Do not label axes only with units. In the example, write “Magnetization (A/m)” or “Magnetization {A[m(1)]}”, not just “A/m”. Do not label axes with a ratio of quantities and units. For example, write “Temperature (K)”, not “Temperature/K”.

II. INTRODUCTION

III. BACKGROUND

A. *k*-nearest Neighbours (*k*NN)

The nearest neighbour algorithm is an example of instance-based learning where the prediction for an instance is made by comparing the instance to similar training instances. This is a lazy learning approach since we delay processing training data until prediction is needed [1]. Since *k*NN is also similarity-based, it is non-parametric, which means that it does not learn a mapping from the input space to the output space.

The *k*-nearest neighbour algorithm is presented in Algorithm 1 which explains that we calculate the distance from each data point in our training dataset to the distance. From there, we choose the closed *k* observations and perform inference. In a classification context we may perform majority

voting to determine the class, in a regression context we use the mean or median of the *k* nearest observations.

In order to use the algorithm we need to define a value of *k* which decides the number of nearest neighbour instances to consider when predicting. Choosing a value a low value for *k* results in a high variance, but low bias leading to an unstable model which tends to overfit. Choosing a large value for *k* results in a smaller variance, but high bias, leading to a more stable model which tending to underfit. Measures such as *k*-fold cross validation assist in the selection of a value of *k*.

Algorithm 1 *k*-Nearest Neighbors Algorithm

```
function KNN( $D, x, k$ )  
  for all  $x' \in D$  do  
     $d = \text{distance}(x, x')$   
  end for  
   $\text{sort}(d)$   
   $S = \text{set of } k \text{ patterns in } D \text{ closest to } x$   
  return class as majority class in  $S$   
end function
```

Notation: D denotes the training dataset, x is the query instance, k is the number of nearest neighbours, and $\text{distance}(\cdot)$ is a distance metric (e.g., Euclidean distance).

Another important part in the *k*NN algorithm is the distance metric. This determines how similarity between observations is measured. Euclidean distance is the most common metric. It calculates the line distance between two points in space, this is useful for continuous numerical features. Manhattan distance computes the sum of absolute differences across each dimension. It is often used when features are not strongly correlated. A generalization of these two is the Minkowski distance, which introduces a parameter p that allows interpolation between different norms: when $p = 1$, it reduces to Manhattan distance, and when $p = 2$, it becomes Euclidean distance. The choice of metric can significantly impact the performance of *k*NN, especially in high-dimensional spaces or when feature scales differ.

B. Classification Trees

Classification trees represent a fundamental approach in supervised machine learning for predictive modelling, where the learned model forms a hierarchical tree structure with non-terminal nodes representing decisions on descriptive features

and terminal leaf nodes representing target feature predictions [2]. A classification tree is a type of decision tree where leaf nodes represent different discrete classes.

Classification trees employ recursive partitioning where a training dataset is systematically divided into increasingly homogeneous subsets based on feature values. It uses a greedy strategy to select the feature to partition upon at each split that maximises information gain [1]. The information gain criterion, derived from Shannon's entropy measure, quantifies the reduction in uncertainty achieved by partitioning the data according to a specific feature test. Formally, given a dataset D with M different classes, the entropy is calculated as:

$$H(D) = - \sum_{m=1}^M p(y_m) \log_M p(y_m)$$

where $p(y_m)$ represents the probability of class y_m occurring in D . When the dataset is partitioned on feature α into O outcomes, the information gain is computed as:

$$\text{gain}(\alpha) = H(D) - H_{\alpha}(D)$$

where

$$H_{\alpha}(D) = \sum_{o=1}^O p_o H(D_o)$$

represents the weighted entropy after the split.

Since we recursively partition until either a stopping condition is satisfied or all subsets are homogeneous, classification trees inherently lead to overfitting, as the algorithm attempts to perfectly classify all training instances, often creating overly complex trees that capture noise rather than underlying patterns [2]. To remedy this, post-pruning techniques are used to remove branches that do not improve generalisation performance.

Classification trees can naturally handle both categorical and numerical features. Classification tree induction algorithms can cope with missing values. Unfortunately, classification trees exhibit an axis-aligned bias which restricts decision boundaries to be parallel to feature axes.

C. Expectations With Respect to Data Quality Issues

Since kNN relies heavily on distance metrics, it will be severely impacted by missing values since it cannot compute meaningful distances between points. This can however be handled by ignoring the missing values in the distance calculation and scaling up the weight of the non-missing descriptive features. The kNN algorithm can actually be used as an imputer itself as well. If nothing is done the algorithm will either fail or give biased distances. Classification trees on the other hand handle missing values more gracefully. Classification trees can adjust the gain ratio calculation

$$\text{gainRatio}(x) = (1 - F) \times (H(D) - H_x(D))$$

where F is the fraction of missing patterns.

With regards to the Facet feature being correlated with the Aspect feature, kNN will be moderately impacted since the distance calculation used will overweight certain areas of similarity. There is no built in mechanism to handle feature correlation as a data issue. Classification trees will be minimally impacted since trees naturally handle correlated features through the feature selection process involving the gain ratio/information gain. This means that the classification tree induction will likely select one feature and ignore the correlated redundant one.

Concerning the Inclination feature only containing noisy values, the impact on the kNN algorithm will depend on the value of k since higher values are more sensitive to noise and lower values are more robust to noise. However, since there are many predictive features the value won't dominate the distance calculation. Classification trees are robust to noise and will handle it through pruning. The noisy patterns will end up in small leaf nodes as the model overfits and post-pruning removes these leaves. The noise becomes a minority and does not affect majority class prediction.

Since kNN employs distance measures which are often sensitive to outliers since the distance can be seriously skewed even if only one feature is outlying. For small values of k outliers are unlikely to be nearest neighbours and will therefore have limited influence. For large values of k the outliers will also have limited influence due to majority voting. Cumulatively, outliers will have a moderate impact on the kNN algorithm. Classification trees are robust to outliers since they handle them in a similar way to noise. Outliers are first isolated in small leaf nodes during induction. Pruning removes leaves which contain outliers which results in outliers becoming a minority in combined larger subsets. This makes decision trees robust to outliers.

The kNN algorithm will be severely impacted from features with numeric ranges that differ significantly from one another since features with large ranges dominate distance calculations. For example, if one feature takes on the value of 0 or 1 and another takes on the value of 0 or 100 000, the second will overwhelm the distance calculation. Without normalisation, the kNN algorithm will be unreliable. Classification trees have a natural robustness to different feature scales and therefore will be minimally impacted. This stems from trees using individual feature values for splitting thresholds.

The mixed data types will affect the kNN algorithm since it requires special handling for mixed types such as using the Gower's similarity with different similarity measures for different types. However, in this dataset since there are two categorical variables with the first encoded over four columns and the second over 40 and 13 numeric features there can be a distance bias toward the numeric features. Euclidean distance assumes features are continuous, but the categorical features are encoded as zero or one. Classification trees naturally handle mixed data types since they directly compare categorical features against each other for splitting, and the splitting is simply decided off information gain.

The impact of the Water Level feature having a cardinality

of one will be none at all for kNN. It will only contribute the the distance calculation unnecessarily and be a computational waste. Since all observations have the same value, all the distance measures will be equally skewed, resulting in no impact at all. The classification tree will never select this feature during tree indication since there is zero information gain because all values are identical. Therefore, classification trees naturally handle this issue and it will have no impact.

The Observation ID unique feature will have a sever impact on kNN since each instance appears maximally different from all others on this feature. The problem is that this dominates distance calculations making all instances seem dissimilar. It breaks the fundamental assumption of similarity-based learning. Classification trees will overfit severely since each unique value would create a separate branch. The information gain maximisation favours features with many outcomes so the result will be an extremely busy tree with poor generalisation.

The skewed class distribution impact on kNN depends on the value of k , since a smaller k is more robust to class imbalance whilst a large k results in the majority class dominating the predictions. Methods such as the Synthetic Minority Over-sampling Technique (SMOTE) must be considered. The impact of classification trees is moderate since they are sensitive to skewed distributions. The minority class instances result in small leaves which are likely to be pruned, then the combined instances will be classification as a majority class. This results in poor minority class recognition.

IV. METHODOLOGY

A. Tools and Libraries

This implementation was created using Python. NumPy [3] was used for numerical computations. Pandas [4] was used for data manipulation and analysis. Matplotlib [5] was used for data visualisation and plotting. Scikit-learn [6] was used extensively for pre-processing, model selection, validation, both machine learning algorithms, and the evaluation metrics. The imbalanced-learn library was used for the SMOTE and Tomek techniques. For statistical testing, SciPy [7] was used. All the random seeds were set to 42 when using these libraries for reproducibility.

V. EMPIRICAL PROCEDURE

A. Data Pre-processing

In this section we discuss the data pre-processing steps applied to the dataset prior to implementing the k-nearest neighbour and classification tree algorithms.

1) *k-Nearest Neighbour Data Pre-processing*: First, the unique identifier feature `Observation_ID` was removed because unique values make all instances appear maximally different. This breaks the similarity-based learning assumptions and leads to the feature dominating distance calculations. The feature has no signal that can provide the model with any benefit.

Second, we removed the constant feature `Water_Level` since it provides zero discriminative information but adds computational overhead to all the distance calculations. We

can be certain that the model performance will not change based on the fact that the feature has a cardinality of one.

Third, we normalised the numeric features by applying linear scaling to ensure the numeric features are in the same range being between zero and one. Features with large ranges will simply dominate distance calculations over smaller range features. This would make the model unreliable and heavily weight features with higher values. Specifically, `Elevation`, `Aspect`, `Slope`, `Horizontal_Distance_To_Hydrology`, `Vertical_Distance_To_Hydrology`, `Horizontal_Distance_To_Roadways`, `Hillshade_9am`, `Hillshade_Noon`, `Hillshade_3pm` are all normalised. The scaler was fit to the training set only in all situations (including the cross fold validation) and applied to the training and test set.

Fourth, we handled the missing values by using kNN imputation with a k of five. There are only 298 observations that are missing the `Slope` feature. For each of these we computed the Euclidean distance between the observation and all others based on available features, then identified the k closest rows (neighbours) that have the missing feature value. The mean of these values were calculated and imputed to the missing observation.

Fifth, we addressed the feature correlation issue by removing the `Facet` feature since it is just providing redundant information. Correlated features can overweight certain aspects in distance calculations resulting in worse model performance. Essentially, the `Aspect` feature would have double the influence on the model as the other features if `Facet` was not removed.

Sixth, the mixed data types issue was handled by converting the one categorical variable `Soil_Type1` to numeric. This was achieved by simply replacing all the instances with a *positive* value with the number one, and the *negative* instances with zero. This approach was chosen since it aligned the feature with the other soil type features which all take on the values of either one or zero.

Seventh, the features containing outliers were left untouched. The kNN algorithm has moderate sensitivity to feature-based outliers, but a small k value provides natural protection.

Eighth, SMOTE was used to improve the class imbalance. Since the five minority classes had significantly fewer observations than the majority two classes, we employed SMOTE to increase the number of minority class observations to 20% of the majority class. This partial balancing approach was used since the dataset only increased in size by 34%, as opposed to the dataset exploding in size if they were all were made equal. For example, if we over-sample class 4 to 283,301 we'd create 280 000 synthetic samples for one class. This is massive and risky for overfitting. On top of this Tomek links for majority class under-sampling was also used to reduce the number of observations in the majority classes. Using this under-sampling technique alone would not be sufficient since only a small fraction of the majority samples are removed. The Tomek link

technique resulted in the majority classes reducing by roughly 60 000 observations in each majority class. That is, Class 1 now has 158,427 samples, and Class 2 now has 225,918 samples after reduction. It is important to note that SMOTE was only applied to the *training* datasets and never to the *test* datasets.

Ninth, the noisy feature *Inclination* was removed since it was confirmed to contain only noise. Noisy features affect distance calculations, although larger k values provide some robustness it is safer to remove it since we know it contains no signal.

2) *Classification Trees Data Pre-processing*: First, the unique identifier feature *Observation_ID* was removed because it creates severe overfitting. Classification tree induction will result in one branch per unique value in this feature, this triggers the many-values bias in information gain maximisation. Again, since the feature provides no predictive power we can safely remove it without losing anything important.

Second, the constant *Water_Level* feature was removed since it will always result in zero information gain if the classification tree were to split on it. The algorithm will naturally ignore it, but it does waste computational resources.

Third, the target class imbalance was addressed using SMOTE to over-sample the minority classes. Again, the minority classes were over-sampled to 20% of the size of the largest majority class. This is done because trees are sensitive to skewed distributions since minority classes get pruned away leading to poor recognition. To use SMOTE, the missing values needed to be temporarily imputed using kNN imputation and the categorical variables encoded. This was performed to enable SMOTE to work as intended. The classification tree was still trained on the dataset containing missing values and categorical features. It is important to note that SMOTE was only applied to the *training* datasets and never to the *test* datasets.

Fourth, the missing values were left as is since the classification trees have excellent built-in missing value handling. Pre-processing is therefore unnecessary in this case. The tree induction algorithm distributes the instance fractionally across all branches where a feature is missing, this is done proportional to the observed frequencies of those branches among the non-missing data.

Fifth, the correlated *Facet* and *Aspect* features are left as is. Trees naturally handle correlation through information gain and redundant features are automatically ignored. The algorithm will select the most informative feature.

Sixth, the mixed data types are left as is. Classification trees natively handle both categorical and numeric features without the need for conversion.

Seventh, the numerical feature scale differences need not be attended to. Trees are scale-invariant since they use a threshold-based split that works regardless of the feature ranges. The features are only compared on an information gain basis and not by their values. This is unlike the kNN algorithm which compares the values of the features during the distance measurements.

Eighth, the outliers are left in the dataset. Classification trees are naturally robust to outliers through the pruning process. The C4.5 algorithm [2] uses error-based pruning. If a branch only exists because of a few outliers, that branch has few training examples. Therefore, its estimated error rate is high.

Ninth, the noisy feature *Inclination* is kept since the pruning process handles noise. Trees handle noise well by isolating noisy patterns in small leaves that get pruned.

B. Control Parameter Tuning

In this section we discuss the custom cross-validation approach used to tune the parameters for each of the algorithms. The key principle is that the SMOTE/Tomek techniques are only applied to training folds, never to validation folds, ensuring realistic performance estimates.

1) *k-Nearest Neighbour Control Parameter Tuning*: The three primary parameters were optimised over during a 5-fold stratified cross-validation on a 20% subset for computational efficiency. First, is the value of k which controls the number of nearest neighbours to consider. Values of [3, 5, 7, 9, 11, 15] were considered. Second, both the options for the weights parameter were considered. Specifically, uniform where all neighbours are weighted equally and distance weighting where closer neighbours have more influence. Third, three distance metrics were considered: Euclidean distance, manhattan distance, and minkowski distance (which is a generalisation of both with a p parameter). Since there are three parameters with six, two, and three options respectively, there are $6 \times 2 \times 3 = 36$ parameter combinations to test in the cross-validation.

2) *Classification Trees Control Parameter Tuning*: The five primary parameters were optimised over during a three-fold stratified cross-validation on a 20% subset for computational efficiency. First, four maximum depth parameter values were considered: [10, 15, 20, 25]. This parameter controls tree depth to balance complexity against generalisation. Second, the minimum number of samples required to split an internal node. The values of [2, 5, 10, 20] were considered. Higher values prevent overfitting by requiring more evidence for splits. Third, the minimum number of samples required at leaf nodes was tuned, with values of [1, 2, 5, 10]. This prevents the formation of overly specific rules that do not generalise well. Fourth, the splitting criterion was evaluated using both *gini* (Gini impurity) and *entropy* (information gain). Gini is computationally faster, while entropy provides a more theoretically principled approach. Finally, the *max_features* parameter was tuned over ['sqrt', 'log2', None], which controls the number of features considered for the best split. Using *sqrt* or *log2* introduces randomness to improve generalisation, while None considers all features. Since there are five parameters with four, four, four, two, and three options respectively, there are $4 \times 4 \times 4 \times 2 \times 3 = 384$ parameter combinations tested.

C. Performance Metrics

Four primary metrics were considered to evaluate the performance of the models. Accuracy was used to evaluate the overall correctness of predictions across all classes. Macro

and weighted precision was used to evaluate the algorithm's ability to avoid false positives. Macro and weighted recall was used to test the algorithm's ability to find all positive cases. F1-Score was also employed as a measure of the algorithm's overall performance. Cohen's Kappa which is a statistic that measures the level of agreement between two classifiers on categorical data and adjusts for the agreement could happen by chance. This metric was employed to compare the predictions of the classifier to the true values of the test set. Similarly, the Matthews Correlation Coefficient (MCC) was employed as a comprehensive performance metric that extends beyond binary classification to evaluate the quality of multiclass predictions, providing a balanced assessment that accounts for the correlation between predicted and actual classifications across all classes while being less sensitive to class imbalance than traditional accuracy measures.

10-fold cross-validation mean accuracy was employed to provide robust performance estimation. The standard deviation of the cross-validation score was used to measure the algorithm stability across the various folds. The range of the cross-validation scores was also evaluated to show the best and worst fold performance.

Statistical testing using t-tests and 95% confidence intervals quantified significance and uncertainty. Computational metrics such as training time, prediction time, and model complexity (e.g., parameter count, tree depth, memory usage) were also recorded. Weighted averages were applied to handle class imbalance, and multiple metrics were used to avoid over-reliance on a single measure.

D. Analysis Process

A custom cross-validation implementation was employed to ensure proper resampling and that the SMOTE and Tomek techniques were only applied within each *training* fold. This ensures the validation was unbiased since it was performed on the original unbalanced data and also prevents data leakage.

The learning curves of both the classification trees and kNN algorithms were compared by sampling increasingly larger training subsets of the data. These training subsets are then used to train the models, obviously in the case of kNN the training data is the data it is *fit* to (since kNN does not train). The accuracy performance was then plotted for both the training and test sets. Each iteration was performed with three-fold cross validation.

While all the tests were run, the number of iterations and running time was recorded. This assisted in the analysis of the computational complexity of the algorithms.

The Mann-Whitney U test [8] was used to perform a pairwise comparison (1×1) between the two models' cross validation results. This is a non-parametric test assumes the data (accuracies in this case) is non-parametric and that the runs of the models are independent. The test can be performed on small samples (in this case 10 per group) making it useful despite having lower power.

For the classification trees, the library used has a built-in feature importance attribute which measures the contribution

of each feature to the reduction of node impurity. This was a valuable tool for understanding which features had the greatest discriminative power in the classification task and provided insights into the underlying data structure.

VI. RESEARCH RESULTS

A. Parameter Tuning Results

1) *K-nearest Neighbours*: The custom five-fold cross-validation grid search returned the best kNN parameters with $k = 3$, the Manhattan distance metric and the distance weight function used in prediction. Table II gives the top five models ordered by mean cross-validation score they achieving during the grid search. The values of k are all low in the table indicating that higher values of k result in the model under-fitting and not performing as well on the validation sets. This table also suggests that the choice of distance metric has a smaller effect on performance compared to the number of neighbours. Additionally, using distance-based weighting appears consistently beneficial, as it prioritises closer neighbours and improves prediction accuracy.

TABLE II
TOP PERFORMING KNN MODELS WITH DIFFERENT HYPERPARAMETERS

| Model | Mean CV Score | Std CV Score | Metric | k | Weights |
|-------|---------------|--------------|-----------|-----|----------|
| 1 | 0.8312 | 0.0019 | manhattan | 3 | distance |
| 2 | 0.8310 | 0.0030 | manhattan | 5 | distance |
| 3 | 0.8295 | 0.0015 | euclidean | 3 | distance |
| 4 | 0.8295 | 0.0015 | minkowski | 3 | distance |
| 5 | 0.8281 | 0.0028 | manhattan | 7 | distance |

Experimentation exhibited in Figure 1 shows how as the value of k increases, the training and test error both decrease and never trend upward again. We confirm that the parameters [3, 5, 7, 9, 11, 15] over which we performed the grid search are reasonable since they should all result in their own balance between underfitting and overfitting. The k value of one was excluded from this search due to the glaring difference which will always occur between the test and the train accuracies. Adding the value would also increase the computational complexity unnecessarily.




Fig. 1. Training vs. test accuracy of kNN algorithm as k increases.

2) *Classification Trees*: The classification trees parameter tuning resulting in the optimal model with the gini criterion being used, a maximum depth of 25, no maximum number of features, the minimum number of samples in a leaf to be one and the minimum number of samples required to split on being two. Table III presents the top five models from the grid search ranked by mean cross-validation score. The results indicate that deeper trees tend to perform better, as shown by the highest scoring model having a depth of 25. Reducing the maximum depth or increasing the minimum number of samples required to split generally leads to a slight decrease in performance, suggesting that the dataset benefits from the ability to model complex interactions.

The small differences in mean cross-validation scores among the top models indicate that the performance is relatively stable across minor variations in parameters. Additionally, the low standard deviations show that the model’s performance is consistent across the folds of cross-validation. Overall, these results suggest that while the decision tree can capture non-linear relationships effectively, careful tuning of depth and split parameters is necessary to balance model complexity and generalisation.

TABLE III
COMPARISON OF DECISION TREE MODELS WITH DIFFERENT
HYPERPARAMETERS

| Model | Mean CV Score | Std CV Score | Parameters |
|-------|---------------|--------------|-----------------------|
| 1 | 0.7965 | 0.0018 | gini, 25, None, 1, 2 |
| 2 | 0.7957 | 0.0037 | gini, 20, None, 1, 5 |
| 3 | 0.7955 | 0.0021 | gini, 25, None, 1, 5 |
| 4 | 0.7949 | 0.0039 | gini, 25, None, 2, 2 |
| 5 | 0.7947 | 0.0026 | gini, 20, None, 1, 10 |

^aParameter order: criterion, max_depth, max_features, min_samples_leaf, min_samples_split.

The effect of the depth parameter is shown by Figure 2. We observe that as the maximum depth of the classification tree increase, the training error decreases as expected. The test error decreases until a point at which it begins to increase again due to overfitting. This is since the model becomes too complex and begins modelling noise as opposed to signal.

B. Cross-validation Results

Note that scores presented in Table III and Table II represent the accuracy on a small subset (20% for kNN and 10% for the classification trees) of the data. The following results are from the model applied to the *entire* dataset.

1) *K-nearest Neighbours*: For each fold in the 10-fold cross-validation, the SMOTE and Tomek techniques were all only applied training samples. This resulted in a kNN mean accuracy of 0.9184 across the 10 folds. The accuracy scores ranged from 0.9170 to 0.9202 which showed that the algorithm generalises well. The high one-vs-rest ROC-AUC score of 0.9721 demonstrates excellent discriminatory ability across all seven classes, significantly outperforming the classification tree model. The model achieves strong precision (0.9196) and recall (0.9184) balance, indicating effective identification

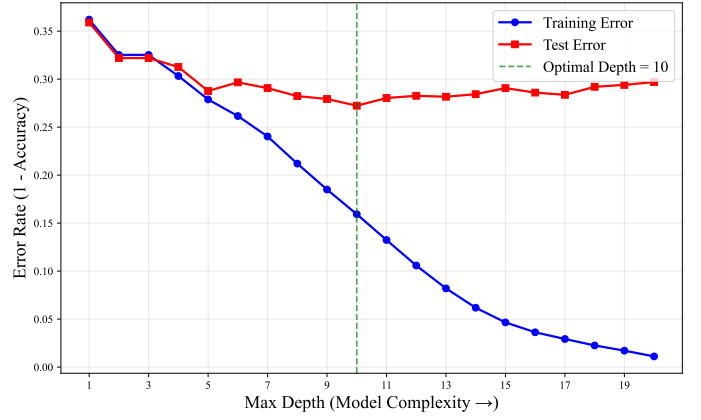


Fig. 2. Effect of maximum depth parameter on classification trees.

across the multiclass problem. The F1-score of 0.9187 confirms the robust balance between precision and recall. The low standard deviations across all metrics indicate stable and reliable classification performance across different data splits. The kNN algorithm also reported a MCC of 0.8591 and a Cohen’s Kappa of 0.8590. These are similar since they are both change-corrected metrics that account for class imbalance. These metrics indicate that we have obtained consistent performance across classes.

TABLE IV
PERFORMANCE METRICS OF KNN AND CLASSIFICATION TREE MODELS
(MEAN \pm STD) OVER 10 FOLDS

| Metric | KNN | Classification Tree |
|-----------|---------------------|---------------------|
| Accuracy | 0.9184 \pm 0.0011 | 0.9276 \pm 0.0023 |
| Precision | 0.9196 \pm 0.0011 | 0.9280 \pm 0.0024 |
| Recall | 0.9184 \pm 0.0011 | 0.9276 \pm 0.0023 |
| F1-Score | 0.9187 \pm 0.0011 | 0.9278 \pm 0.0023 |
| ROC-AUC | 0.9721 \pm 0.0005 | 0.9562 \pm 0.0015 |

2) *Classification Trees*: For the classification trees 10-fold cross validation was performed with SMOTE only applied to the training folds. The maximum accuracy of all runs was 0.9309 and the minimum was 0.9233. The mean accuracy of 0.9276 and other metrics can be observed in Table IV. The very low standard deviation of 0.0023 on the accuracy metric indicates that the model has stable performance. While achieving higher overall accuracy than kNN, the classification tree model shows a lower one-vs-rest ROC-AUC score of 0.9562, suggesting slightly reduced discriminatory power when distinguishing between individual classes and all others. The model demonstrates excellent precision (0.9280), recall (0.9276), and F1-score (0.9278) performance. The superior performance across accuracy, precision, recall, and F1-score metrics suggests better overall classification effectiveness compared to the kNN approach. The consistently low standard deviations across all metrics confirm the model’s stability and reliability in multiclass prediction tasks. The classification trees algorithm reported a MCC of 0.8196 and a Cohen’s Kappa of 0.8196. These are similar since they are both change-

corrected metrics that account for class imbalance. We observe that these values indicate excellent classification performance, with the model achieving strong agreement beyond what would be expected by random chance and demonstrating robust predictive capability even in the presence of imbalanced class distributions.

C. Learning Curves

Figure 4 provides a visual illustration of how the classification tree learns the signal in the data as the size of the training set increases. This figure was generated using only 10% of the data for computational purposes. We observe the upward trend of the validation score and the slight downward trend of the training score. This is indicative of the model improving on its generalisation performance. Considering these trends, the results reported in Table IV make sense. We see how the validation score would eventually increase to approximately 0.9160 as the training set increases in size to the full dataset. The learning curve for the kNN algorithm in Figure 3 is similar, however due to the small sample dataset, the training score is near perfect. This will change as more minority classes are included in the larger dataset.

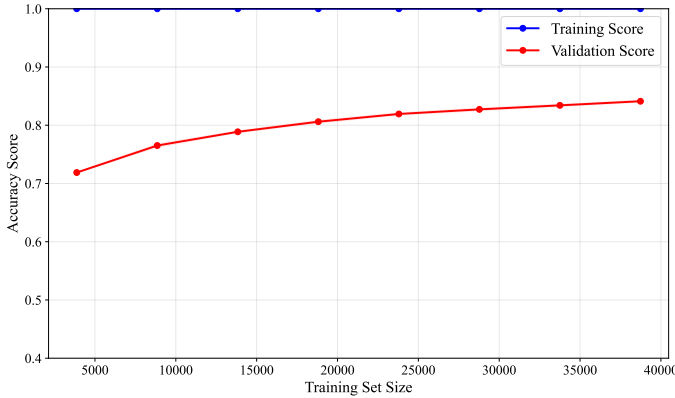


Fig. 3. The learning curve of the kNN algorithm.

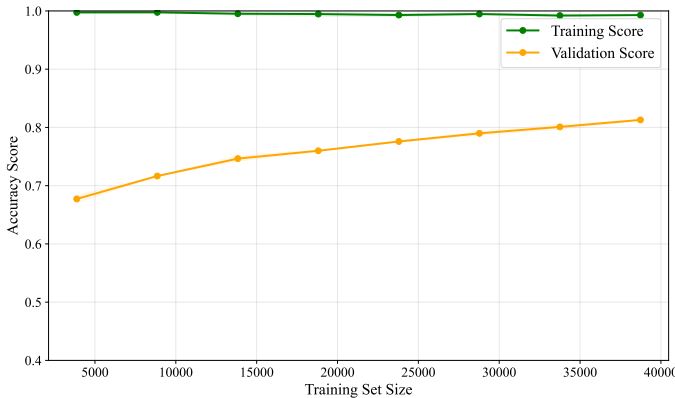


Fig. 4. The learning curve of the classification trees algorithm.

D. Computational Analysis

The majority of the computational complexity stemmed from the grid search approach used in the hyper parameters tuning. Despite sampling a significantly smaller portion of the data, the number of parameters combinations (36 and 384 respectively) still increased the computation time drastically. Since five-fold cross validation was used for each combination in kNN, there were $36 \times 5 = 180$ fits to run. Only three folds were used for decision trees due to the significantly larger number of combinations. This resulted in $384 \times 3 = 1152$ fits to run. The addition of the Tomek link undersampling greatly increased the running time of the parameter tuning process. This is due to the SMOTE and Tomek (for kNN) undersampling being applied only to the training data in *each* iteration of the parameter tuning process. One complex optimisation would be to store the calculations required for this technique and re-use them over each iteration, however this would require a substantial amount of memory and care to prevent data leakage. The Tomek link procedure adds considerable computational overhead as it requires identifying and removing borderline instances through nearest neighbor calculations for every data point in each training fold.

To evaluate the performance of the models thoroughly, it was imperative that all the data was used. Thus, 10-fold cross validation over the entire dataset was employed to ensure robust performance estimates. The computational analysis revealed significant differences in training efficiency, with each fold requiring approximately 12.6 minutes for kNN and 2.1 minutes for Decision Trees, highlighting the substantially higher computational cost associated with kNN's distance-based lazy learning approach compared to the tree-based eager learning methodology.

E. Statistical Significance Testing

The Mann-Whitney U test [8] was performed using the 10 cross validation scores previously described. This resulted in a p-value of 0.00018 which is very small, therefore indicating there is strong evidence to reject the null hypothesis that the classification trees and the kNN scores come from the same distribution.

F. Feature Importance

The most useful features for reducing the impurity of the nodes in the classification tree algorithm are seen in Table V. The feature `Elevation` contributes the most to the reducing the Gini impurity. Note that the absolute magnitude of these feature importance values have no intrinsic meaning, since they are normalised to sum to one. Interpreting the relative magnitude of the `Elevation` feature indicates that it contributes roughly three times as much as the second most important feature, `Horizontal_Distance_To_Roadways`. A shortfall of this table is since the many soil type features are encoded with dummy variables, they are not individually as informative on their own, but collectively may contribute a even larger amount to the reduction of impurity.

TABLE V
TOP 10 FEATURE IMPORTANCES FROM DECISION TREE CLASSIFIER

| Feature | Importance |
|------------------------------------|------------|
| Elevation | 0.3463 |
| Horizontal_Distance_To_Roadways | 0.1246 |
| Horizontal_Distance_To_Fire_Points | 0.1074 |
| Horizontal_Distance_To_Hydrology | 0.0607 |
| Vertical_Distance_To_Hydrology | 0.0314 |
| Soil_Type3 | 0.0289 |
| Wilderness_Area3 | 0.0277 |
| Hillshade_9am | 0.0248 |
| Hillshade_Noon | 0.0243 |
| Wilderness_Area1 | 0.0190 |

VII. CONCLUSION

REFERENCES

- [1] J. Kelleher, B. Namee, and A. D'Arcy, *Fundamentals of Machine Learning for Predictive Data Analytics, second edition: Algorithms, Worked Examples, and Case Studies*. MIT Press, 2020. [Online]. Available: <https://books.google.co.za/books?id=wtGMEAAQBAJ>
- [2] J. R. Quinlan, *C4.5: Programs for Machine Learning*. San Francisco, CA, USA: Morgan Kaufmann Publishers Inc., 1993.
- [3] C. R. Harris, K. J. Millman, S. J. van der Walt, R. Gommers, P. Virtanen, D. Cournapeau, E. Wieser, J. Taylor, S. Berg, N. J. Smith, R. Kern, M. Picus, S. Hoyer, M. H. van Kerkwijk, M. Brett, A. Haldane, J. F. del Río, M. Wiebe, P. Peterson, P. Gérard-Marchant, K. Sheppard, T. Reddy, W. Weckesser, H. Abbasi, C. Gohlke, and T. E. Oliphant, "Array programming with NumPy," *Nature*, vol. 585, no. 7825, pp. 357–362, Sep. 2020. [Online]. Available: <https://doi.org/10.1038/s41586-020-2649-2>
- [4] T. pandas development team, "pandas-dev/pandas: Pandas," Feb. 2020. [Online]. Available: <https://doi.org/10.5281/zenodo.3509134>
- [5] J. D. Hunter, "Matplotlib: A 2d graphics environment," *Computing in Science & Engineering*, vol. 9, no. 3, pp. 90–95, 2007.
- [6] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay, "Scikit-learn: Machine learning in Python," *Journal of Machine Learning Research*, vol. 12, pp. 2825–2830, 2011.
- [7] P. Virtanen, R. Gommers, T. E. Oliphant, M. Haberland, T. Reddy, D. Cournapeau, E. Burovski, P. Peterson, W. Weckesser, J. Bright, S. J. van der Walt, M. Brett, J. Wilson, K. J. Millman, N. Mayorov, A. R. J. Nelson, E. Jones, R. Kern, E. Larson, C. J. Carey, Í. Polat, Y. Feng, E. W. Moore, J. VanderPlas, D. Laxalde, J. Perktold, R. Cimrman, I. Henriksen, E. A. Quintero, C. R. Harris, A. M. Archibald, A. H. Ribeiro, F. Pedregosa, P. van Mulbregt, and SciPy 1.0 Contributors, "SciPy 1.0: Fundamental Algorithms for Scientific Computing in Python," *Nature Methods*, vol. 17, pp. 261–272, 2020.
- [8] H. B. Mann, "On a test of whether one of two random variables is stochastically larger than the other," *The annals of mathematical statistics*, vol. 18, no. 1, pp. 50–60, 1947.